

Intro to Data Analysis with R

Pri Oberoi

5/16/2016

While you are waiting, go here to get started:

www.github.com/prioberoi/R_intro_to_data_analysis




Pri Oberoi (poberoi@doc.gov)
Data Scientist, Commerce Data Service
US Department of Commerce

Goals

Our goals for the class

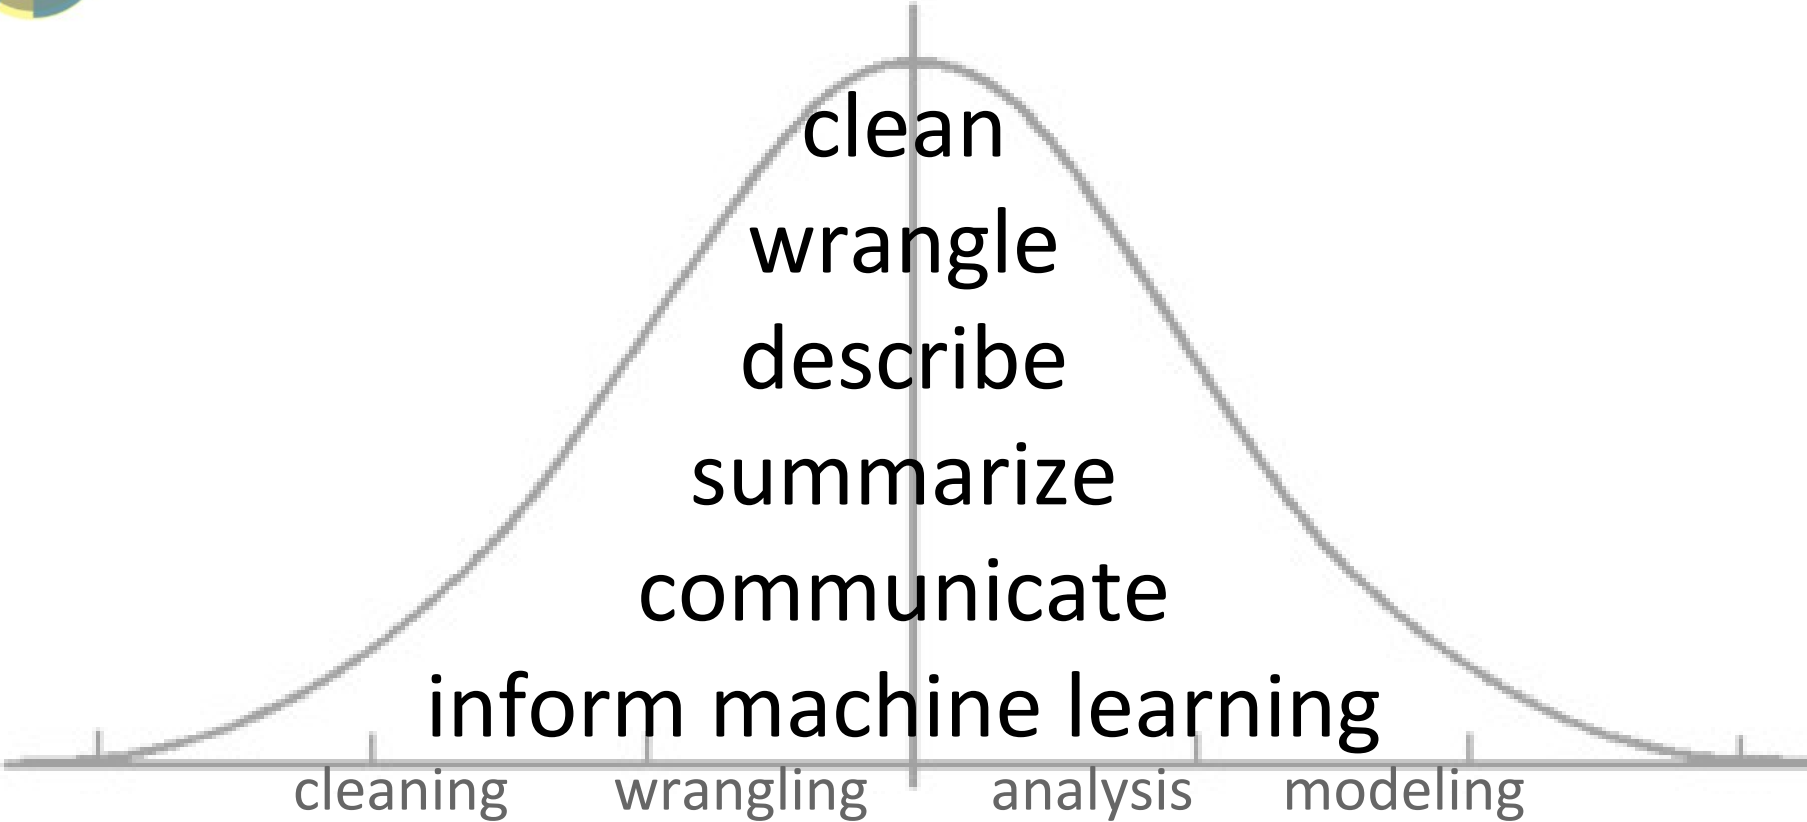
- Import, clean, visualize data in R
- Role of data analysis in the data science pipeline
- Use data analysis to summarize,
-

The Data Science Pipeline



Why do data analysis?

clean
wrangle
describe
summarize
communicate
inform machine learning



clean
wrangle
describe
summarize
communicate
inform machine learning


cleaning

wrangling

analysis

modeling

The Pipeline



R Markdown

Output formats:

HTML, PDF, MS Word, HTML5 slides,
books, dashboards, websites

Benefits:

Easy to create

Embedded R code chunks

(which can be visible or not on the final output)

Allows you to add a narrative through your code

Reproducible

PUT SCREENSHOT OF R MARKDOWN FOR
CLASS HERE.
ANNOTATE.

Analysis Toolkit

Visualization

Statistics

Aggregation

Data Visualization

ggplot2

qplot()

“quick plot”

- similar to plot() from base R
- less typing
- less customizable

```
qplot(carat, price, data = diamonds,  
size = I(1), alpha = I(1/10), main =  
"qplot scatter plot")
```

ggplot()

- more customizable
- more functionality

```
ggplot(data = diamonds, aes(x = carat,  
y = price)) +  
  geom_point(size = 1, alpha = 1/10) +  
  ggtitle("ggplot scatter plot")
```

ggplot2

x

y

```
qplot(carat, price, data = diamonds,  
size = I(1), alpha = I(1/10), main =  
"qplot scatter plot")
```

aesthetics like
point size, point
transparency,
plot title

data

y

```
ggplot(data = diamonds, aes(x = carat,  
y = price)) +  
geom_point(size = 1, alpha = 1/10) +  
ggtitle("ggplot scatter plot")
```

x

Your turn

Run the code in chunk 2: **Scatterplots**

Update `ggplot()` code so the color of the scatterplot points varies based on the value of 'cut'

You can do this by adding a 'colour =' argument to the `aes()` mapping

Your turn

Run the code in chunk 3: **Histograms and Bar Charts**

Note that you can set the 'binwidth' for histograms

Run the code in chunk 4: **Boxplots and Violin Plots**

Box plots: more widely interpretable

Violin plots: useful for non-normal distributions and to scale to
number of observations

NTIA Broadband Data Example

NTIA's broadband data from June, 2014 for Washington, DC

Hypothesis Testing

Null hypothesis: the typical upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds

Your turn

Import the data by running chunk 5

Look at the dataframe

```
View(data)
dim(data)
names(data)
str(data)
summary(data)
```

Create a histogram of max advertised download speeds
(maxaddown)

10 min break

Cleaning

Messy Data

Signs you have messy data:

- Column headers are values, not variable names
- Multiple variables are stored in one column
- Variables are stored in both rows and columns
- Multiple types of observational units are stored in the same table
- A single observational unit is stored in multiple tables



[This content is from Hadley Wickham's paper on tidy data](#)

Column headers are values, not variable names

We will be looking at the maxaddown, maxadup, typicdown, typicup variables

They are stored as different columns/variables, rather than different values.

```
# Run chunk 6
```

```
melt() # this is a function that converts columns into rows  
?melt # use ?melt to read the documentation on this function
```


Multiple variables are stored in one column

`data_clean` now has one column named 'variable' that contains the variable indicating if this is advertised or typical as well as whether this speed is for uploads or downloads.

```
# Run chunk 7 and look at the resulting data_clean dataframe
```

```
data_clean$speedDirection <- "download"  
data_clean$speedDirection[data_clean$variable %in% c("maxadup","typicup")]  
<- "upload"  
data_clean$speedSource <- "advertised"  
data_clean$speedSource[data_clean$variable %in% c("typicup","typicdown")]  
<- "typical"  
data_clean$variable <- NULL
```

Variables are stored in both rows and columns

We don't have this problem in our dataset, but here is an example:

id	year	month	element	d1	d2	d3	d4	d5	d6	d7	d8
MX17004	2010	1	tmax	—	—	—	—	—	—	—	—
MX17004	2010	1	tmin	—	—	—	—	—	—	—	—
MX17004	2010	2	tmax	—	27.3	24.1	—	—	—	—	—
MX17004	2010	2	tmin	—	14.4	14.4	—	—	—	—	—
MX17004	2010	3	tmax	—	—	—	—	32.1	—	—	—
MX17004	2010	3	tmin	—	—	—	—	14.2	—	—	—
MX17004	2010	4	tmax	—	—	—	—	—	—	—	—
MX17004	2010	4	tmin	—	—	—	—	—	—	—	—
MX17004	2010	5	tmax	—	—	—	—	—	—	—	—
MX17004	2010	5	tmin	—	—	—	—	—	—	—	—

Table 11: Original weather dataset. There is a column for each possible day in the month. Columns d9 to d31 have been omitted to conserve space.

Variables are individual columns (id, year, month), spread across columns (day, d1–d31) and across rows (tmin, tmax)

Multiple types of observational units are stored in the same table

data_clean has data on different observational units, the provider/holding company, broadband speeds, location

Repeating values in a column are a result.

Similar to database normalization.

Note, some data analysis tools work with denormalized data

A single observational unit is stored in multiple tables


Data values for a single variable are found across tables

```
# the following functions are helpful  
rbind() #add dataframe as rows, must have same number of columns  
cbind() #add dataframe as columns, must have same number of rows  
merge() #merge two data frames by common columns or row names  
ldply() #from the plyr package reads multiple csvs into one dataframe
```

Summarizing Data Within Groups

Summarizing Data Within Groups

This boxplot indicates advertised and typical speeds differ.



```
# aggregate() will run functions over a group you specify
```

```
# what does this do:
```

```
# data_clean[data_clean$speedDirection == 'download','value']
```

```
aggregate(data_clean[data_clean$speedDirection == 'download','value'], list  
(data_clean[data_clean$speedDirection == 'download', 'speedSource']), mean)
```

Your turn

Update chunk 8 to show the mean upload speeds by
speedSource (advertised or typical)

10 min break


Correlation

Correlation

We expect that the advertised and typical speeds are correlated


Run chunk 9

Correlation



Pearson's product-moment correlation

```
data: data$maxaddown and data$typicdown  
t = 242.74, df = 149810, p-value < 2.2e-16  
alternative hypothesis: true correlation is not equal to 0  
95 percent confidence interval:  
 0.5276687 0.5349375  
sample estimates:  
      cor  
0.5313129
```




p-value < 0.05
not surprising

Revisit our null hypothesis

Null hypothesis: the typical upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds

Comparing Samples

Null hypothesis: the typical upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds



Null hypothesis: the typical download speeds for broadband providers in Washington, DC are the same as the advertised speeds

Let's do a quick t-test to see if that difference is statistically significant

```
t.test(response ~ variable, data)
```

Null hypothesis: the typical download speeds for broadband providers in Washington, DC are the same as the advertised speeds

Let's do a quick t-test to see if that difference is statistically significant

Two hand-drawn grey arrows. One starts from the word 'difference' in the text above and points to the 'response' variable in the code. The other starts from the word 'statistically' and points to the 'data' argument in the code.

```
t.test(response ~ variable, data)
```

```
#does the download speed differ based on whether the speed was  
#advertised or typical in our dataset?
```


Null hypothesis: the typical mean upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds

H_0

Run chunk 10

```
> t.test(data_clean[data_clean$speedDirection == 'download', 'value'] ~ data_clean[data_clean$speedDirection == 'download', 'speedSource'], data_clean)
```

Welch Two Sample t-test

```
data: data_clean[data_clean$speedDirection == "download", "value"] by data_clean[data_clean$speedDirection == "download", "speedSource"]
t = 266.22, df = 243900, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 2.031413 2.061546
sample estimates:
mean in group advertised      mean in group typical
      5.346951              3.300471
```

P-value < 0.05

H_a

Your turn!

update the t-test code (chunk 10) to see if the typical upload speed is different than the advertised upload speed

10 min break

From statistical tests to statistical learning

Can we predict what the typical upload and download speeds are if given the advertised upload and download speeds?

Linear Regression

```
fit <- lm(response ~ predictors, data)
```

Linear Regression

```
# run chunk 11  
fit <- lm(typicdown ~ maxaddown, data)
```

```
Call:  
lm(formula = typicdown ~ maxaddown, data = data)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.778	-2.096	1.116	1.798	2.222

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.479633	0.020473	-72.27	<2e-16 ***
maxaddown	0.893987	0.003683	242.74	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.167 on 149806 degrees of freedom
Multiple R-squared: 0.2823, Adjusted R-squared: 0.2823
F-statistic: 5.892e+04 on 1 and 149806 DF, p-value: < 2.2e-16


our predictors and
their significance

residuals: how far off were the
predicted values from the observed
value?

p-value < 0.05

hypothetical typical
download speed for an
advertised download
speed of 0

Linear Regression



Resources

R-bloggers: <http://www.r-bloggers.com/>

FlowingData: <http://flowingdata.com/category/tutorials/>

Google's R Style Guide: <https://google.github.io/styleguide/Rguide.xml>

R Markdown: <http://rmarkdown.rstudio.com/>